

Ming Hill
CS 660
Professor Athanassoulis
28 Nov 2023

PA3 Write Up Worked with Abdelazim

For our Filter class, there weren't many design decisions that we had to make. For the smaller methods (open, close, rewind, etc) we just called the parent operator function before calling on the fields themselves. This is also true when we implemented those functions for the other classes in this assignment. Other than that, for the `fetchNext()` method, we decided that the simplest way to approach this problem was to just iterate through the child iterator, until we found a tuple that passed the filter, and then returned it.

For the join class, one decision that was pivotal in completing the `fetchNext()` method was to create a `t1` field. The `t1` field represented an optional tuple object, which would store the current tuple from `child1`. For the `fetchNext()` method, we used a nested loop, matching a given `t1` to a `t2`. We added this line "`t1 != std::nullopt || child1->hasNext()`" for the conditional of the outer loop, which would make sure that we would always have a `child1`, and if not, we would just return a `nullopt`. Then by iterating through `child2`, we can find the element that would match with `t1`, and return the tuple. When `child2` is at the end of the iteration, we know then to iterate `child1` to the next value therefore, we would rewind `child2` back to the beginning and iterate `child1` one iteration. With this logic, we would be able to match and join each element from both groups. This is the same logic that we used for `HashEquiJoin`.

For our `IntegerAggregate`, a couple of small design decisions were made to optimize the implementation of the methods. These included adding a `hashMaps` field to the `IntegerAggregate` class (`count`, `groupSum`), which we could refer to in `mergeTupletoGroup`. The `count` map represents the number of counts per grouping, if there is a grouping, and `groupSum` represents the total value of each grouping. For `mergeTupleIntoGroup`, we first identify if there is a grouping or not, and set that as our `groupbyField` variable. Then by using that field variable, we can increase the count in our count hashmap of the respective grouping if there is one. Then by using case matching on the aggregate operators, we implemented the correct operations for updating `groupSum`. For `next()` in `IntegerAggregatorIterator`, we get the key from the `open()` method, which points us to the current key in the count hashmap. This key is the same for both count and `groupSum` hashmaps, therefore we can retrieve the respective values for both hashmaps. Then once again, similar to `mergeTupletoGroup`, we utilize case matching on the aggregate operator to decide which operation to use and store the result in the `resultValue` variable. Finally, depending on whether there is a grouping or not, we create a tuple with one field, which is just the aggregate value, or a tuple with

two fields, one being the grouping and the second being the value. Finally, we increment our current iterator to point at the next key in the count hashmap before returning our newly created tuple.

For our Aggregate class, the `fetchNext()` method is related to the `IntegerAggregator` class by fetching the next tuple and returning it. Additionally, `getTupleDesc` is dependent on the grouping field. Depending on whether there is a grouping field or not, the `TupleDesc` will contain a group and aggregate column, or only an aggregate column.

For the Delete/Insert classes, our implementation for their respective `fetchNext()` is very similar. We just iterate through the child and keep track of the number of tuples that we insert/delete, before returning a new tuple that that value as the first field.

We did not make any significant changes to the API, the only thing that we changed was inserting an unordered map to the constructor of `agregateIterator`.

We did not miss or have any incomplete elements of code.

This assignment was less intense and time-consuming than PA2 for my partner and me. Many of the earlier methods we had to implement were pretty straightforward (Filter, Insert, Delete). We had 2 main issues when trying to complete the code. The first one was the `fetchNext()` method for the join operator. When we initially tried developing the method, we had a major logic mistake, instead of matching up each element from the first group (`t1`) to every matching element in the second group (`t2`), we would completely iterate over `t1` after the first match. The second issue we had was with the `IntegerAggregate` file. Given the minimal instructions, it was difficult at first to understand what we were supposed to implement, and when we figured that out, the logic behind the implementation took a while to figure out. Specifically, we were confused about how the `IntegerAggregatorIterator` was tied to `IntegerAggregate`, as well as how to implement `next()` in the iterator class and `fetchNext()` in the `IntegerAggregate` class.

For this project, I worked with Abdel from our class. For 90% of the project, we worked together, especially on more of the difficult problems such as `IntegerAggregator` and `fetchNext()` for join. Other than that, we worked separately on the last two classes, Delete and Insert, however, after implementing both of them, we went over them together to get an understanding of the logic behind our implementations.